

6 Topic modeling

In text mining, we often have collections of documents, such as blog posts or news articles, that we'd like to divide into natural groups so that we can understand them separately. Topic modeling is a method for unsupervised classification of such documents, similar to clustering on numeric data, which finds natural groups of items even when we're not sure what we're looking for.

Latent Dirichlet allocation (LDA) is a particularly popular method for fitting a topic model. It treats each document as a mixture of topics, and each topic as a mixture of words. This allows documents to “overlap” each other in terms of content, rather than being separated into discrete groups, in a way that mirrors typical use of natural language.

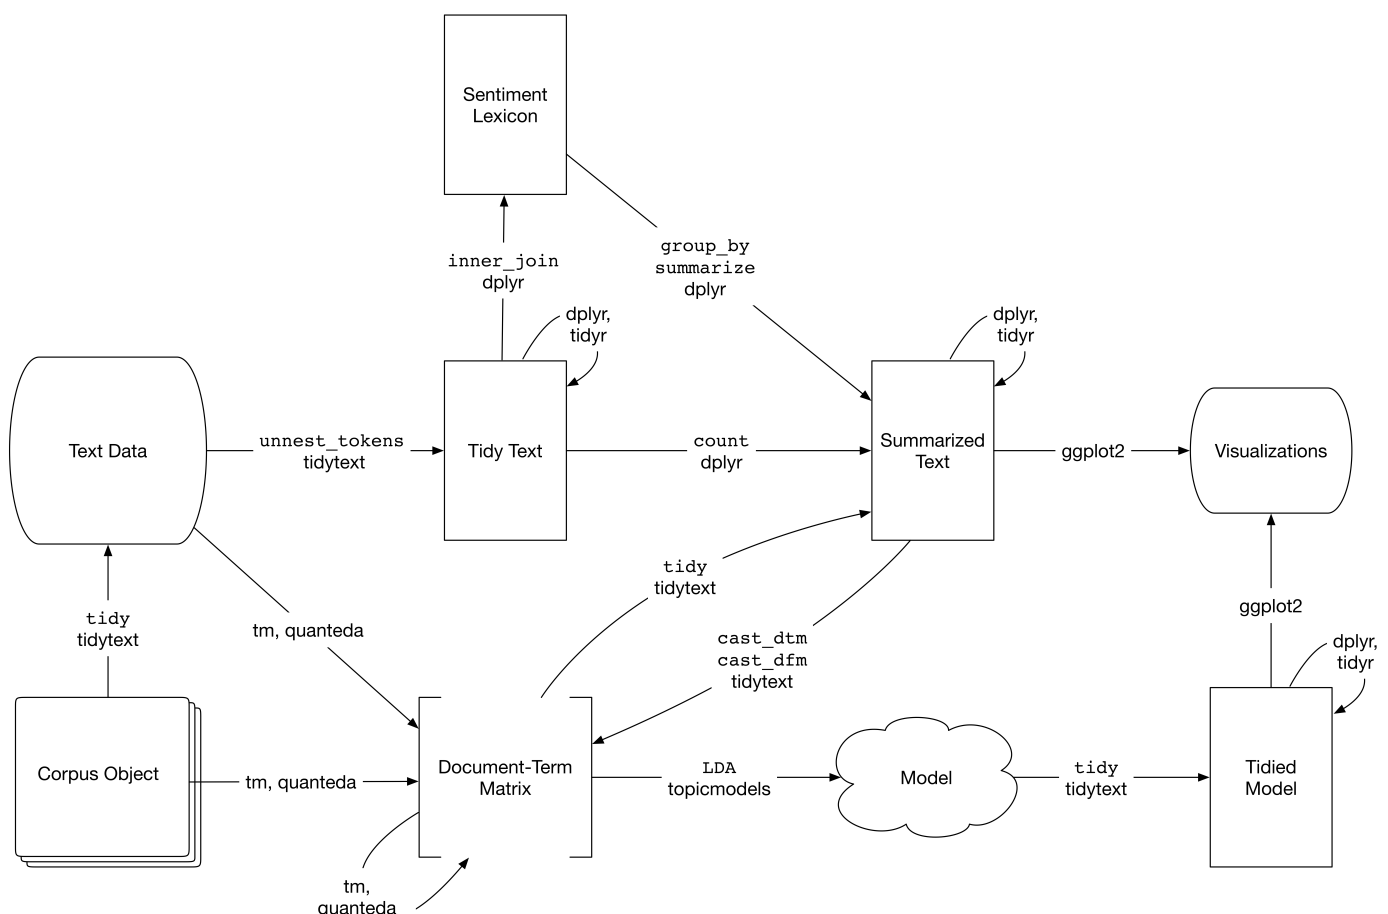


Figure 6.1: A flowchart of a text analysis that incorporates topic modeling. The `topicmodels` package takes a Document-Term Matrix as input and produces a model that can be tidied by `tidytext`, such that it can be manipulated and visualized with `dplyr` and `ggplot2`.

As Figure 6.1 shows, we can use tidy text principles to approach topic modeling with the same set of tidy tools we've used throughout this book. In this chapter, we'll learn to work with LDA objects from the `topicmodels` package, particularly tidying such models so that they can be manipulated with `ggplot2` and `dplyr`. We'll also explore an example of clustering chapters from several books, where we can see that a topic model “learns” to tell the difference between the four books based on the text content.

6.1 Latent Dirichlet allocation

Latent Dirichlet allocation is one of the most common algorithms for topic modeling. Without diving into the math behind the model, we can understand it as being guided by two principles.

- **Every document is a mixture of topics.** We imagine that each document may contain words from several topics in particular proportions. For example, in a two-topic model we could say “Document 1 is 90% topic A and 10% topic B, while Document 2 is 30% topic A and 70% topic B.”
- **Every topic is a mixture of words.** For example, we could imagine a two-topic model of American news, with one topic for “politics” and one for “entertainment.” The most common words in the politics topic might be “President”, “Congress”, and “government”, while the entertainment topic may be made up of words such as “movies”, “television”, and “actor”. Importantly, words can be shared between topics; a word like “budget” might appear in both equally.

LDA is a mathematical method for estimating both of these at the same time: finding the mixture of words that is associated with each topic, while also determining the mixture of topics that describes each document. There are a number of existing implementations of this algorithm, and we'll explore one of them in depth.

In Chapter 5 we briefly introduced the `AssociatedPress` dataset provided by the `topicmodels` package, as an example of a `DocumentTermMatrix`. This is a collection of 2246 news articles from an American news agency, mostly published around 1988.

```
library(topicmodels)
```

```
data("AssociatedPress")
```

```
AssociatedPress
```

```
## <<DocumentTermMatrix (documents: 2246, terms: 10473)>>
## Non-/sparse entries: 302031/23220327
## Sparsity           : 99%
## Maximal term length: 18
## Weighting          : term frequency (tf)
```

We can use the `LDA()` function from the `topicmodels` package, setting `k = 2`, to create a two-topic LDA model.

Almost any topic model in practice will use a larger `k`, but we will soon see that this analysis approach extends to a larger number of topics.

This function returns an object containing the full details of the model fit, such as how words are associated with topics and how topics are associated with documents.

```
# set a seed so that the output of the model is predictable
ap_lda <- LDA(AssociatedPress, k = 2, control = list(seed = 1234))
ap_lda

## A LDA_VEM topic model with 2 topics.
```

Fitting the model was the “easy part”: the rest of the analysis will involve exploring and interpreting the model using tidying functions from the `tidytext` package.

6.1.1 Word-topic probabilities

In Chapter 5 we introduced the `tidy()` method, originally from the `broom` package (Robinson 2017), for tidying model objects. The `tidytext` package provides this method for extracting the per-topic-per-word probabilities, called β (“beta”), from the model.

```
library(tidytext)

ap_topics <- tidy(ap_lda, matrix = "beta")
ap_topics
```

```
## # A tibble: 20,946 x 3
##   topic term      beta
##   <int> <chr>    <dbl>
## 1     1 1 aaron  1.69e-12
## 2     2 2 aaron  3.90e- 5
## 3     3 1 abandon 2.65e- 5
## 4     4 2 abandon 3.99e- 5
## 5     5 1 abandoned 1.39e- 4
## 6     6 2 abandoned 5.88e- 5
## 7     7 1 abandoning 2.45e-33
## 8     8 2 abandoning 2.34e- 5
## 9     9 1 abbott  2.13e- 6
## 10    10 2 abbott  2.97e- 5
## # ... with 20,936 more rows
```

Notice that this has turned the model into a one-topic-per-term-per-row format. For each combination, the model computes the probability of that term being generated from that topic. For example, the term “aaron” has a 1.686917×10^{-12} probability of being generated from topic 1, but a 3.8959408×10^{-5} probability of being generated from topic 2.

We could use `dplyr`’s `top_n()` to find the 10 terms that are most common within each topic. As a tidy data frame, this lends itself well to a `ggplot2` visualization (Figure 6.2).

```

library(ggplot2)
library(dplyr)

ap_top_terms <- ap_topics %>%
  group_by(topic) %>%
  top_n(10, beta) %>%
  ungroup() %>%
  arrange(topic, -beta)

ap_top_terms %>%
  mutate(term = reorder_within(term, beta, topic)) %>%
  ggplot(aes(term, beta, fill = factor(topic))) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ topic, scales = "free") +
  coord_flip() +
  scale_x_reordered()

```

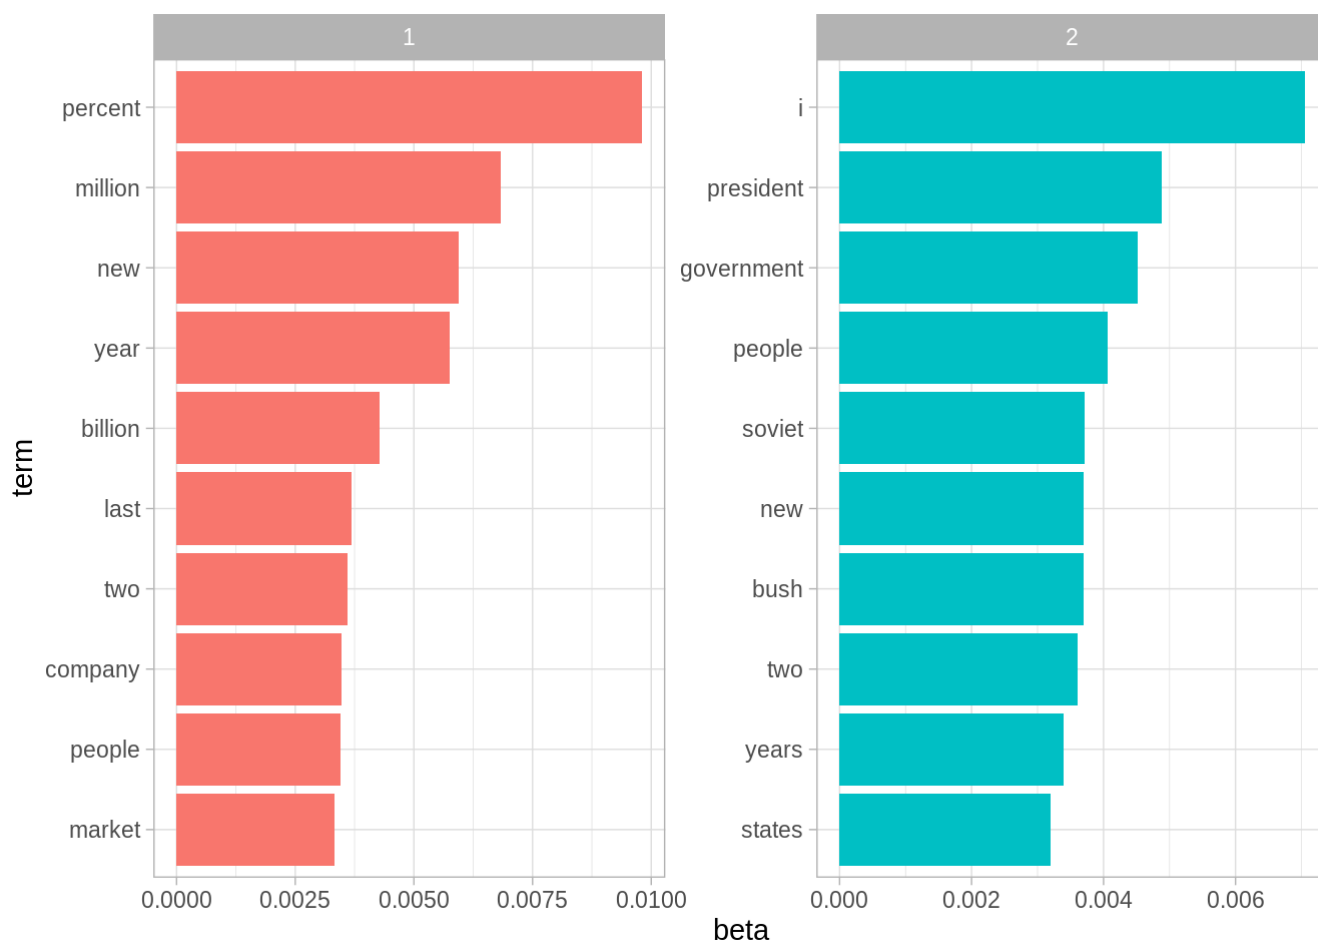


Figure 6.2: The terms that are most common within each topic

This visualization lets us understand the two topics that were extracted from the articles. The most common words in topic 1 include “percent”, “million”, “billion”, and “company”, which suggests it may represent business or financial news. Those most common in topic 2 include “president”, “government”, and “soviet”, suggesting that this topic represents political news. One important observation about the words in each topic is that some words, such as “new” and “people”, are common within both topics. This is an advantage of topic modeling as opposed to “hard clustering” methods: topics used in natural language could have some overlap in terms of words.

As an alternative, we could consider the terms that had the *greatest difference* in β between topic 1 and topic 2. This can be estimated based on the log ratio of the two: $\log_2(\frac{\beta_2}{\beta_1})$ (a log ratio is useful because it makes the difference symmetrical: β_2 being twice as large leads to a log ratio of 1, while β_1 being twice as large results in -1). To constrain it to a set of especially relevant words, we can filter for relatively common words, such as those that have a β greater than 1/1000 in at least one topic.

```
library(tidyr)

beta_spread <- ap_topics %>%
  mutate(topic = paste0("topic", topic)) %>%
  spread(topic, beta) %>%
  filter(topic1 > .001 | topic2 > .001) %>%
  mutate(log_ratio = log2(topic2 / topic1))

beta_spread
```

```
## # A tibble: 198 x 4
##   term          topic1    topic2 log_ratio
##   <chr>         <dbl>    <dbl>    <dbl>
## 1 administration 0.000431 0.00138      1.68
## 2 ago            0.00107 0.000842    -0.339
## 3 agreement      0.000671 0.00104      0.630
## 4 aid            0.0000476 0.00105      4.46
## 5 air           0.00214 0.000297    -2.85
## 6 american       0.00203 0.00168    -0.270
## 7 analysts       0.00109 0.000000578 -10.9
## 8 area          0.00137 0.000231    -2.57
## 9 army          0.000262 0.00105      2.00
## 10 asked         0.000189 0.00156      3.05
## # ... with 188 more rows
```

The words with the greatest differences between the two topics are visualized in Figure 6.3.




Figure 6.3: Words with the greatest difference in β between topic 2 and topic 1

We can see that the words more common in topic 2 include political parties such as “democratic” and “republican”, as well as politician’s names such as “dukakis” and “gorbachev”. Topic 1 was more characterized by currencies like “yen” and “dollar”, as well as financial terms such as “index”, “prices” and “rates”. This helps confirm that the two topics the algorithm identified were political and financial news.

6.1.2 Document-topic probabilities

Besides estimating each topic as a mixture of words, LDA also models each document as a mixture of topics. We can examine the per-document-per-topic probabilities, called γ (“gamma”), with the `matrix = "gamma"` argument to `tidy()`.

```
ap_documents <- tidy(ap_lda, matrix = "gamma")
ap_documents
```

```
## # A tibble: 4,492 x 3
##   document topic    gamma
##   <int> <int>    <dbl>
## 1      1      1  0.248
## 2      2      2  0.362
## 3      3      3  0.527
## 4      4      4  0.357
## 5      5      5  0.181
## 6      6      6  0.000588
## 7      7      7  0.773
## 8      8      8  0.00445
## 9      9      9  0.967
## 10     10     10  0.147
## # ... with 4,482 more rows
```

Each of these values is an estimated proportion of words from that document that are generated from that topic. For example, the model estimates that only about 25% of the words in document 1 were generated from topic 1.

We can see that many of these documents were drawn from a mix of the two topics, but that document 6 was drawn almost entirely from topic 2, having a γ from topic 1 close to zero. To check this answer, we could `tidy()` the document-term matrix (see Chapter 5.1) and check

what the most common words in that document were.

```
tidy(AssociatedPress) %>%
  filter(document == 6) %>%
  arrange(desc(count))
```

```
## # A tibble: 287 x 3
##   document term      count
##   <int> <chr>    <dbl>
## 1      6 noriega      16
## 2      6 panama      12
## 3      6 jackson       6
## 4      6 powell        6
## 5      6 administration  5
## 6      6 economic       5
## 7      6 general        5
## 8      6 i              5
## 9      6 panamanian      5
## 10     6 american        4
## # ... with 277 more rows
```

Based on the most common words, this appears to be an article about the relationship between the American government and Panamanian dictator Manuel Noriega, which means the algorithm was right to place it in topic 2 (as political/national news).

6.2 Example: the great library heist

When examining a statistical method, it can be useful to try it on a very simple case where you know the “right answer”. For example, we could collect a set of documents that definitely relate to four separate topics, then perform topic modeling to see whether the algorithm can correctly distinguish the four groups. This lets us double-check that the method is useful, and gain a sense of how and when it can go wrong. We’ll try this with some data from classic literature.

Suppose a vandal has broken into your study and torn apart four of your books:

- *Great Expectations* by Charles Dickens

- *The War of the Worlds* by H.G. Wells
- *Twenty Thousand Leagues Under the Sea* by Jules Verne
- *Pride and Prejudice* by Jane Austen

This vandal has torn the books into individual chapters, and left them in one large pile. How can we restore these disorganized chapters to their original books? This is a challenging problem since the individual chapters are **unlabeled**: we don't know what words might distinguish them into groups. We'll thus use topic modeling to discover how chapters cluster into distinct topics, each of them (presumably) representing one of the books.

We'll retrieve the text of these four books using the `gutenbergr` package introduced in Chapter 3.

```
titles <- c("Twenty Thousand Leagues under the Sea", "The War of the Worlds",  
           "Pride and Prejudice", "Great Expectations")
```

```
library(gutenbergr)
```

```
books <- gutenberg_works(title %in% titles) %>%  
  gutenberg_download(meta_fields = "title")
```

As pre-processing, we divide these into chapters, use `tidytext`'s `unnest_tokens()` to separate them into words, then remove `stop_words`. We're treating every chapter as a separate "document", each with a name like `Great Expectations_1` or `Pride and Prejudice_11`. (In other applications, each document might be one newspaper article, or one blog post).

```

library(stringr)

# divide into documents, each representing one chapter
by_chapter <- books %>%
  group_by(title) %>%
  mutate(chapter = cumsum(str_detect(text, regex("^chapter ", ignore_case = TRUE)))) %>%
  ungroup() %>%
  filter(chapter > 0) %>%
  unite(document, title, chapter)

# split into words
by_chapter_word <- by_chapter %>%
  unnest_tokens(word, text)

# find document-word counts
word_counts <- by_chapter_word %>%
  anti_join(stop_words) %>%
  count(document, word, sort = TRUE) %>%
  ungroup()

word_counts

```

```

## # A tibble: 104,721 x 3
##   document          word      n
##   <chr>            <chr>  <int>
## 1 Great Expectations_57 joe      88
## 2 Great Expectations_7  joe     70
## 3 Great Expectations_17 biddy    63
## 4 Great Expectations_27 joe     58
## 5 Great Expectations_38 estella  58
## 6 Great Expectations_2  joe     56
## 7 Great Expectations_23 pocket   53
## 8 Great Expectations_15 joe     50
## 9 Great Expectations_18 joe     50
## 10 The War of the Worlds_16 brother  50
## # ... with 104,711 more rows

```

6.2.1 LDA on chapters

Right now our data frame `word_counts` is in a tidy form, with one-term-per-document-per-row, but the `topicmodels` package requires a `DocumentTermMatrix`. As described in Chapter 5.2, we can cast a one-token-per-row table into a `DocumentTermMatrix` with `tidytext`'s `cast_dtm()`.

```
chapters_dtm <- word_counts %>%  
  cast_dtm(document, word, n)  
  
chapters_dtm  
  
## <<DocumentTermMatrix (documents: 193, terms: 18215)>>  
## Non-/sparse entries: 104721/3410774  
## Sparsity           : 97%  
## Maximal term length: 19  
## Weighting          : term frequency (tf)
```

We can then use the `LDA()` function to create a four-topic model. In this case we know we're looking for four topics because there are four books; in other problems we may need to try a few different values of `k`.

```
chapters_lda <- LDA(chapters_dtm, k = 4, control = list(seed = 1234))  
chapters_lda  
  
## A LDA_VEM topic model with 4 topics.
```

Much as we did on the Associated Press data, we can examine per-topic-per-word probabilities.

```
chapter_topics <- tidy(chapters_lda, matrix = "beta")  
chapter_topics
```

```
## # A tibble: 72,860 x 3
##   topic term      beta
##   <int> <chr>    <dbl>
## 1     1   1 joe    5.83e-17
## 2     2   2 joe    3.19e-57
## 3     3   3 joe    4.16e-24
## 4     4   4 joe    1.45e- 2
## 5     5   1 biddy   7.85e-27
## 6     6   2 biddy   4.67e-69
## 7     7   3 biddy   2.26e-46
## 8     8   4 biddy   4.77e- 3
## 9     9   1 estella 3.83e- 6
## 10    10   2 estella 5.32e-65
## # ... with 72,850 more rows
```

Notice that this has turned the model into a one-topic-per-term-per-row format. For each combination, the model computes the probability of that term being generated from that topic. For example, the term “joe” has an almost zero probability of being generated from topics 1, 2, or 3, but it makes up 1% of topic 4.

We could use `dplyr`’s `top_n()` to find the top 5 terms within each topic.

```
top_terms <- chapter_topics %>%
  group_by(topic) %>%
  top_n(5, beta) %>%
  ungroup() %>%
  arrange(topic, -beta)

top_terms
```

```
## # A tibble: 20 x 3
##   topic term      beta
##   <int> <chr>    <dbl>
## 1     1 1 elizabeth 0.0141
## 2     1 1 darcy    0.00881
## 3     1 1 miss     0.00871
## 4     1 1 benet    0.00695
## 5     1 1 jane     0.00650
## 6     2 2 captain  0.0155
## 7     2 2 nautilus 0.0131
## 8     2 2 sea      0.00885
## 9     2 2 nemo     0.00871
## 10    2 2 ned      0.00803
## 11    3 3 people   0.00680
## 12    3 3 martians 0.00651
## 13    3 3 time     0.00535
## 14    3 3 black    0.00528
## 15    3 3 night    0.00448
## 16    4 4 joe      0.0145
## 17    4 4 time     0.00685
## 18    4 4 pip      0.00682
## 19    4 4 looked   0.00637
## 20    4 4 miss     0.00623
```

This tidy output lends itself well to a ggplot2 visualization (Figure 6.4).

```
library(ggplot2)
```

```
top_terms %>%
  mutate(term = reorder_within(term, beta, topic)) %>%
  ggplot(aes(term, beta, fill = factor(topic))) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ topic, scales = "free") +
  coord_flip() +
  scale_x_reordered()
```



Figure 6.4: The terms that are most common within each topic

These topics are pretty clearly associated with the four books! There's no question that the topic of "captain", "nautilus", "sea", and "nemo" belongs to *Twenty Thousand Leagues Under the Sea*, and that "jane", "darcy", and "elizabeth" belongs to *Pride and Prejudice*. We see "pip" and "joe" from *Great Expectations* and "martians", "black", and "night" from *The War of the Worlds*. We also notice that, in line with LDA being a "fuzzy clustering" method, there can be words in common between multiple topics, such as "miss" in topics 1 and 4, and "time" in topics 3 and 4.

6.2.2 Per-document classification

Each document in this analysis represented a single chapter. Thus, we may want to know which topics are associated with each document. Can we put the chapters back together in the correct books? We can find this by examining the per-document-per-topic probabilities, γ ("gamma").

```
chapters_gamma <- tidy(chapters_lda, matrix = "gamma")
```

```
chapters_gamma
```

```
## # A tibble: 772 x 3
##   document          topic    gamma
##   <chr>            <int>    <dbl>
## 1 Great Expectations_57      1 0.0000135
## 2 Great Expectations_7      1 0.0000147
## 3 Great Expectations_17     1 0.0000212
## 4 Great Expectations_27     1 0.0000192
## 5 Great Expectations_38     1 0.354
## 6 Great Expectations_2      1 0.0000172
## 7 Great Expectations_23     1 0.551
## 8 Great Expectations_15     1 0.0168
## 9 Great Expectations_18     1 0.0000127
## 10 The War of the Worlds_16  1 0.0000108
## # ... with 762 more rows
```

Each of these values is an estimated proportion of words from that document that are generated from that topic. For example, the model estimates that each word in the Great Expectations_57 document has only a 0% probability of coming from topic 1 (Pride and Prejudice).

Now that we have these topic probabilities, we can see how well our unsupervised learning did at distinguishing the four books. We'd expect that chapters within a book would be found to be mostly (or entirely), generated from the corresponding topic.

First we re-separate the document name into title and chapter, after which we can visualize the per-document-per-topic probability for each (Figure 6.5).

```
chapters_gamma <- chapters_gamma %>%
```

```
  separate(document, c("title", "chapter"), sep = "_", convert = TRUE)
```

```
chapters_gamma
```



```
## # A tibble: 772 x 4
##   title                chapter topic      gamma
##   <chr>                <int> <int>    <dbl>
## 1 Great Expectations     57     1 0.0000135
## 2 Great Expectations      7     1 0.0000147
## 3 Great Expectations    17     1 0.0000212
## 4 Great Expectations    27     1 0.0000192
## 5 Great Expectations    38     1 0.354
## 6 Great Expectations      2     1 0.0000172
## 7 Great Expectations    23     1 0.551
## 8 Great Expectations    15     1 0.0168
## 9 Great Expectations    18     1 0.0000127
## 10 The War of the Worlds  16     1 0.0000108
## # ... with 762 more rows
```

reorder titles in order of topic 1, topic 2, etc before plotting

```
chapters_gamma %>%
  mutate(title = reorder(title, gamma * topic)) %>%
  ggplot(aes(factor(topic), gamma)) +
  geom_boxplot() +
  facet_wrap(~ title)
```

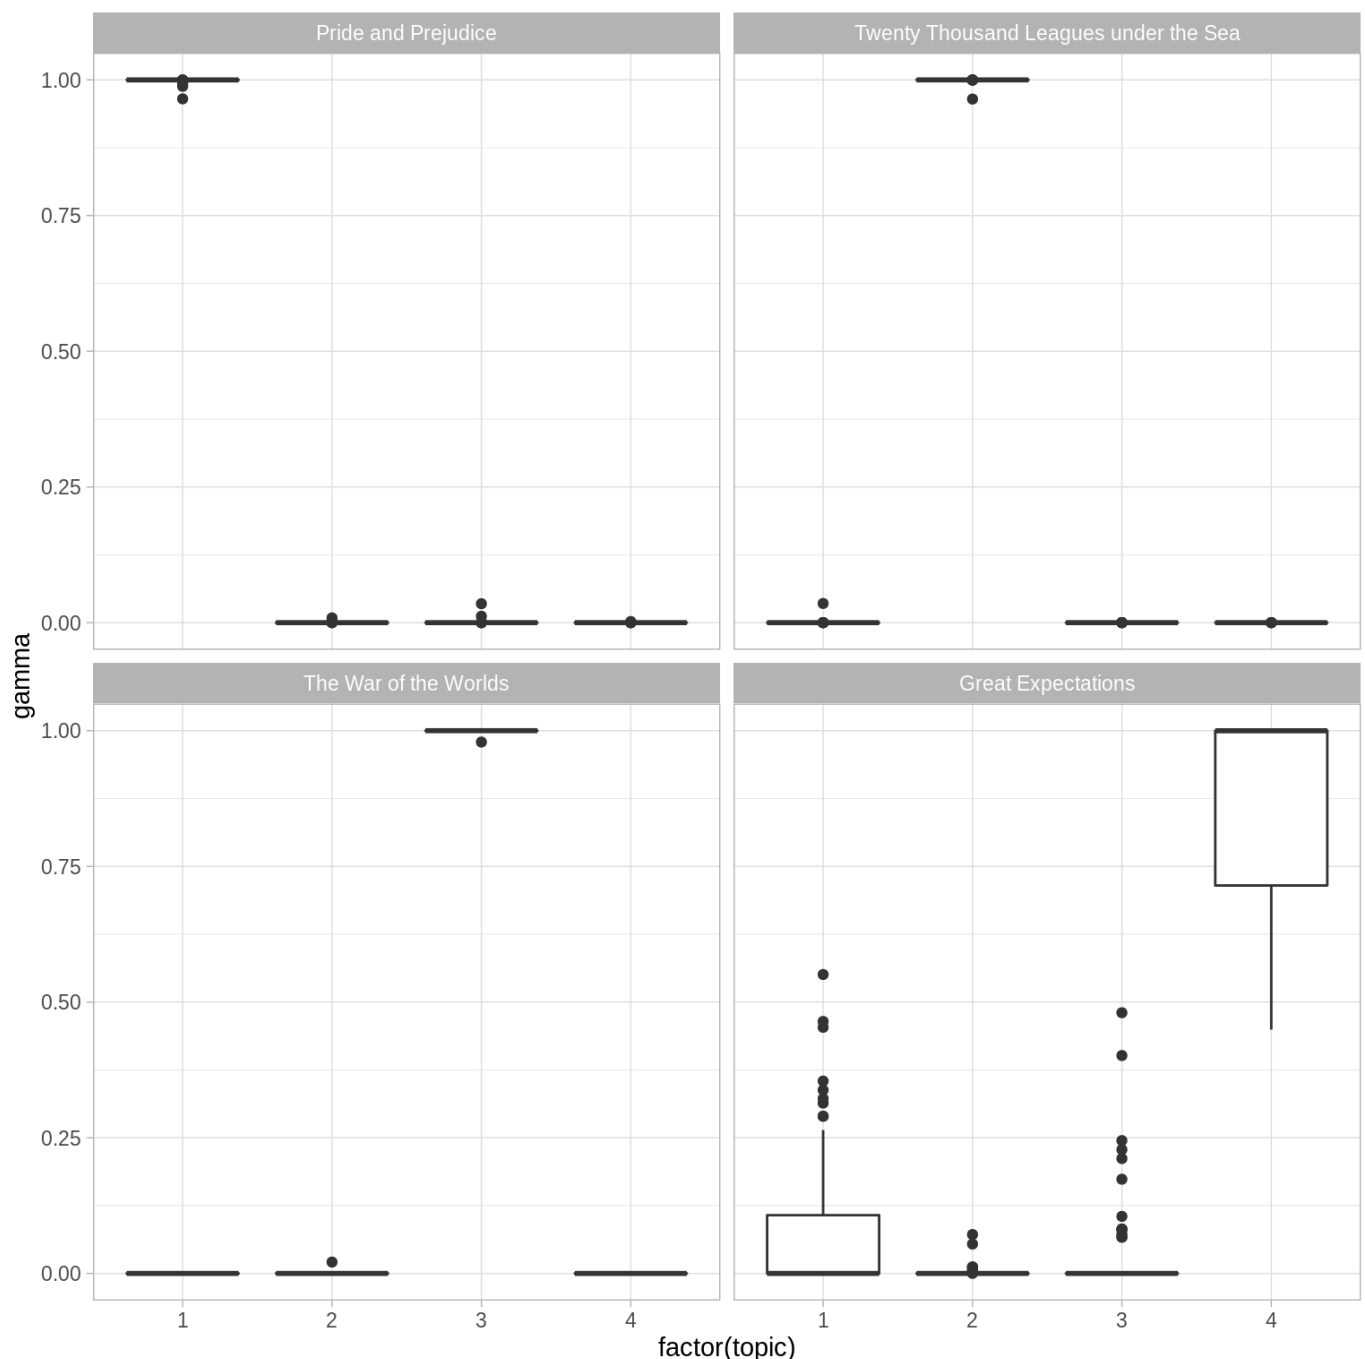


Figure 6.5: The gamma probabilities for each chapter within each book

We notice that almost all of the chapters from *Pride and Prejudice*, *War of the Worlds*, and *Twenty Thousand Leagues Under the Sea* were uniquely identified as a single topic each.

It does look like some chapters from *Great Expectations* (which should be topic 4) were somewhat associated with other topics. Are there any cases where the topic most associated with a chapter belonged to another book? First we'd find the topic that was most associated with each chapter using `top_n()`, which is effectively the "classification" of that chapter.

```
chapter_classifications <- chapters_gamma %>%
  group_by(title, chapter) %>%
  top_n(1, gamma) %>%
  ungroup()
```

```
chapter_classifications
```

```
## # A tibble: 193 x 4
##   title                chapter topic gamma
##   <chr>                <int> <int> <dbl>
## 1 Great Expectations    23     1 0.551
## 2 Pride and Prejudice    43     1 1.000
## 3 Pride and Prejudice    18     1 1.000
## 4 Pride and Prejudice    45     1 1.000
## 5 Pride and Prejudice    16     1 1.000
## 6 Pride and Prejudice    29     1 1.000
## 7 Pride and Prejudice    10     1 1.000
## 8 Pride and Prejudice     8     1 1.000
## 9 Pride and Prejudice    56     1 1.000
## 10 Pride and Prejudice   47     1 1.000
## # ... with 183 more rows
```

We can then compare each to the “consensus” topic for each book (the most common topic among its chapters), and see which were most often misidentified.

```
book_topics <- chapter_classifications %>%
  count(title, topic) %>%
  group_by(title) %>%
  top_n(1, n) %>%
  ungroup() %>%
  transmute(consensus = title, topic)

chapter_classifications %>%
  inner_join(book_topics, by = "topic") %>%
  filter(title != consensus)
```

```
## # A tibble: 2 x 5
##   title                chapter topic gamma consensus
##   <chr>                <int> <int> <dbl> <chr>
## 1 Great Expectations    23     1 0.551 Pride and Prejudice
## 2 Great Expectations    54     3 0.480 The War of the Worlds
```

We see that only two chapters from *Great Expectations* were misclassified, as LDA described one as coming from the “Pride and Prejudice” topic (topic 1) and one from The War of the Worlds (topic 3). That’s not bad for unsupervised clustering!

6.2.3 By word assignments: `augment`

One step of the LDA algorithm is assigning each word in each document to a topic. The more words in a document are assigned to that topic, generally, the more weight (`gamma`) will go on that document-topic classification.

We may want to take the original document-word pairs and find which words in each document were assigned to which topic. This is the job of the `augment()` function, which also originated in the `broom` package as a way of tidying model output. While `tidy()` retrieves the statistical components of the model, `augment()` uses a model to add information to each observation in the original data.

```
assignments <- augment(chapters_lda, data = chapters_dtm)
assignments
```

```
## # A tibble: 104,721 x 4
##   document      term count .topic
##   <chr>        <chr> <dbl> <dbl>
## 1 Great Expectations_57 joe      88     4
## 2 Great Expectations_7  joe      70     4
## 3 Great Expectations_17 joe        5     4
## 4 Great Expectations_27 joe      58     4
## 5 Great Expectations_2  joe      56     4
## 6 Great Expectations_23 joe        1     4
## 7 Great Expectations_15 joe     50     4
## 8 Great Expectations_18 joe     50     4
## 9 Great Expectations_9  joe     44     4
## 10 Great Expectations_13 joe     40     4
## # ... with 104,711 more rows
```

This returns a tidy data frame of book-term counts, but adds an extra column: `.topic`, with the topic each term was assigned to within each document. (Extra columns added by `augment` always start with `.`, to prevent overwriting existing columns). We can combine this `assignments` table with the consensus book titles to find which words were incorrectly classified.

```
assignments <- assignments %>%
  separate(document, c("title", "chapter"), sep = "_", convert = TRUE) %>%
  inner_join(book_topics, by = c(".topic" = "topic"))

assignments
```

```
## # A tibble: 104,721 x 6
##   title                chapter term  count .topic consensus
##   <chr>                <int> <chr> <dbl>  <dbl> <chr>
## 1 Great Expectations      57 joe    88     4 Great Expectations
## 2 Great Expectations       7 joe    70     4 Great Expectations
## 3 Great Expectations     17 joe     5     4 Great Expectations
## 4 Great Expectations     27 joe    58     4 Great Expectations
## 5 Great Expectations       2 joe    56     4 Great Expectations
## 6 Great Expectations     23 joe     1     4 Great Expectations
## 7 Great Expectations     15 joe    50     4 Great Expectations
## 8 Great Expectations     18 joe    50     4 Great Expectations
## 9 Great Expectations      9 joe    44     4 Great Expectations
## 10 Great Expectations    13 joe    40     4 Great Expectations
## # ... with 104,711 more rows
```

This combination of the true book (`title`) and the book assigned to it (`consensus`) is useful for further exploration. We can, for example, visualize a **confusion matrix**, showing how often words from one book were assigned to another, using `dplyr`'s `count()` and `ggplot2`'s `geom_tile` (Figure 6.6).

```
library(scales)
```

```
assignments %>%
  count(title, consensus, wt = count) %>%
  group_by(title) %>%
  mutate(percent = n / sum(n)) %>%
  ggplot(aes(consensus, title, fill = percent)) +
  geom_tile() +
  scale_fill_gradient2(high = "red", label = percent_format()) +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1),
        panel.grid = element_blank()) +
  labs(x = "Book words were assigned to",
       y = "Book words came from",
       fill = "% of assignments")
```




Figure 6.6: Confusion matrix showing where LDA assigned the words from each book. Each row of this table represents the true book each word came from, and each column represents what book it was assigned to.

We notice that almost all the words for *Pride and Prejudice*, *Twenty Thousand Leagues Under the Sea*, and *War of the Worlds* were correctly assigned, while *Great Expectations* had a fair number of misassigned words (which, as we saw above, led to two chapters getting misclassified).

What were the most commonly mistaken words?

```
wrong_words <- assignments %>%
  filter(title != consensus)

wrong_words
```

```
## # A tibble: 4,535 x 6
##   title                                chapter term      count .topic
##   <chr>                                <int> <chr>    <dbl>  <dbl>
## 1 Great Expectations                  38 brother     2      1
## 2 Great Expectations                  22 brother     4      1
## 3 Great Expectations                  23 miss        2      1
## 4 Great Expectations                  22 miss       23      1
## 5 Twenty Thousand Leagues under the Sea    8 miss        1      1
## 6 Great Expectations                  31 miss        1      1
## 7 Great Expectations                   5 sergeant   37      1
## 8 Great Expectations                  46 captain     1      2
## 9 Great Expectations                  32 captain     1      2
## 10 The War of the Worlds                17 captain     5      2
##   consensus
##   <chr>
## 1 Pride and Prejudice
## 2 Pride and Prejudice
## 3 Pride and Prejudice
## 4 Pride and Prejudice
## 5 Pride and Prejudice
## 6 Pride and Prejudice
## 7 Pride and Prejudice
## 8 Twenty Thousand Leagues under the Sea
## 9 Twenty Thousand Leagues under the Sea
## 10 Twenty Thousand Leagues under the Sea
## # ... with 4,525 more rows
```

```
wrong_words %>%
  count(title, consensus, term, wt = count) %>%
  ungroup() %>%
  arrange(desc(n))
```



```
## # A tibble: 3,500 x 4
##   title                consensus      term      n
##   <chr>                <chr>      <chr>    <dbl>
## 1 Great Expectations Pride and Prejudice love      44
## 2 Great Expectations Pride and Prejudice sergeant  37
## 3 Great Expectations Pride and Prejudice lady      32
## 4 Great Expectations Pride and Prejudice miss      26
## 5 Great Expectations The War of the Worlds boat      25
## 6 Great Expectations Pride and Prejudice father    19
## 7 Great Expectations The War of the Worlds water    19
## 8 Great Expectations Pride and Prejudice baby      18
## 9 Great Expectations Pride and Prejudice flopson   18
## 10 Great Expectations Pride and Prejudice family    16
## # ... with 3,490 more rows
```

We can see that a number of words were often assigned to the Pride and Prejudice or War of the Worlds cluster even when they appeared in Great Expectations. For some of these words, such as “love” and “lady”, that’s because they’re more common in Pride and Prejudice (we could confirm that by examining the counts).

On the other hand, there are a few wrongly classified words that never appeared in the novel they were misassigned to. For example, we can confirm “flopson” appears only in *Great Expectations*, even though it’s assigned to the “Pride and Prejudice” cluster.

```
word_counts %>%
  filter(word == "flopson")
```

```
## # A tibble: 3 x 3
##   document                word      n
##   <chr>                <chr>  <int>
## 1 Great Expectations_22 flopson    10
## 2 Great Expectations_23 flopson     7
## 3 Great Expectations_33 flopson     1
```

The LDA algorithm is stochastic, and it can accidentally land on a topic that spans multiple books.

6.3 Alternative LDA implementations

The `LDA()` function in the `topicmodels` package is only one implementation of the latent Dirichlet allocation algorithm. For example, the `mallet` package (Mimno 2013) implements a wrapper around the `MALLET` Java package for text classification tools, and the `tidytext` package provides tidiers for this model output as well.

The `mallet` package takes a somewhat different approach to the input format. For instance, it takes non-tokenized documents and performs the tokenization itself, and requires a separate file of stopwords. This means we have to collapse the text into one string for each document before performing LDA.

```
library(mallet)

# create a vector with one string per chapter
collapsed <- by_chapter_word %>%
  anti_join(stop_words, by = "word") %>%
  mutate(word = str_replace(word, "'", "")) %>%
  group_by(document) %>%
  summarize(text = paste(word, collapse = " "))

# create an empty file of "stopwords"
file.create(empty_file <- tempfile())

docs <- mallet.import(collapsed$document, collapsed$text, empty_file)

mallet_model <- MalletLDA(num.topics = 4)
mallet_model$loadDocuments(docs)
mallet_model$train(100)
```

Once the model is created, however, we can use the `tidy()` and `augment()` functions described in the rest of the chapter in an almost identical way. This includes extracting the probabilities of words within each topic or topics within each document.

```
# word-topic pairs
tidy(mallet_model)

# document-topic pairs
tidy(mallet_model, matrix = "gamma")

# column needs to be named "term" for "augment"
term_counts <- rename(word_counts, term = word)
augment(mallet_model, term_counts)
```

We could use ggplot2 to explore and visualize the model in the same way we did the LDA output.

6.4 Summary

This chapter introduces topic modeling for finding clusters of words that characterize a set of documents, and shows how the `tidy()` verb lets us explore and understand these models using `dplyr` and `ggplot2`. This is one of the advantages of the tidy approach to model exploration: the challenges of different output formats are handled by the tidying functions, and we can explore model results using a standard set of tools. In particular, we saw that topic modeling is able to separate and distinguish chapters from four separate books, and explored the limitations of the model by finding words and chapters that it assigned incorrectly.

References

Mimno, David. 2013. *mallet: A Wrapper Around the Java Machine Learning Tool Mallet*. <https://cran.r-project.org/package=mallet>.

Robinson, David. 2017. *broom: Convert Statistical Analysis Objects into Tidy Data Frames*. <https://CRAN.R-project.org/package=broom>.